

Manoj Ram Mopati

manojrammopati11@gmail.com · 07810294727 · GitHub · LinkedIn · London, UK

Profile

Data Scientist with 4+ years of hands-on experience in data manipulation and data modelling, statistical analysis and generation of insights across retail, consumer and insurance domains using SQL, Python, R, Power BI and Tableau. Proficient in building end-to-end machine learning models and ensuring model deployment through MLOps practices. Over the past year, complemented this experience by building end-to-end GenAI projects during an MSc in Data Science, including LLM-powered applications, RAG systems and NLP solutions, while demonstrating a strong commitment to continuous learning, independent problem-solving and adaptability to changing business needs.

Core Skills

Programming: Python, PySpark, SQL, R.

Analytics: Microsoft Power BI, Power Platforms, Power Query, DAX, Excel, Tableau, Looker Studio.

Statistical & Business Modelling: Marketing Mix Modelling (MMM), regression analysis, statistical modelling, hypothesis testing, A/B testing, time-series forecasting, RFM segmentation, causal analysis, experimentation.

Data Science: Pandas, NumPy, Scikit-learn, SciPy, TensorFlow, PyTorch, data quality checks, data cleaning, data wrangling, data pre-processing, data visualisation, feature engineering, feature selection, log transformation, hyperparameter tuning, one-hot encoding, model deployment, model evaluation, model lifecycle management.

ML & AI: Supervised learning (regression, classification), tree-based and ensemble models, unsupervised learning (clustering), time-series forecasting, neural networks, NLP, RFM segmentation, LLMs, Generative AI, RAG, vector embeddings, semantic search, agentic AI, prompt engineering.

Cloud and MLOps: Azure Databricks, Data Factory, Data Lake Storage, Azure SQL, Azure Machine Learning, Azure Synapse Analytics, Azure Key Vault, Microsoft Entra ID, managed identities, service principals, Azure RBAC, Azure DevOps, Git/GitHub, CI/CD.

Other Tools: HDFS fundamentals, MapReduce concepts, data privacy, GDPR awareness, REST APIs, ETL, KPIs, data warehousing, vector databases, adhoc analysis, Apache Airflow, Agile, Jira, Jupyter Notebooks, R Studio, VS Code.

Work Experience

Junior Data Analyst | Tahir Group | UK | Sep 2025 – Present

Part-time: Sep 2025 – May 2026 | Full-time: May 2026 – Present

- Analysed store-level sales, transaction and labour data using SQL and Python, identifying revenue trends, performance gaps and labour-efficiency opportunities across QSR operations.
- Developed Power BI dashboards and KPI reporting covering revenue, transactions, average transaction value, labour cost %, sales per labour hour and store productivity to support operational decision-making.
- Applied statistical analysis and Difference-in-Differences methodology to evaluate operational interventions, comparing pre/post performance across treatment and comparison stores to quantify incremental business impact.

Data Scientist | Infosys Limited - Client : Retail | India | Sep 2023 – May 2025

- Developed a Power BI analytics platform for a \$1B+ U.S. confectionery portfolio, improving profitable revenue management performance.
- Built Marketing Mix Models (MMM) to quantify marketing and trade promotion effectiveness, enabling data-driven investment optimization.
- Processed and analyzed 3+ years of sales, retail, and promotional records using SQL, Python, PySpark and R to support predictive modelling and commercial analytics.
- Designed trade promotion optimization solutions that improved trade spend efficiency by 8% and supported profit-maximizing promotional strategies.
- Delivered actionable insights on promotional ROI, sales attribution, and revenue drivers to support strategic commercial decision-making.
- Built and deployed scalable analytics workflows using Azure Databricks, Data Factory, Data Lake, Azure SQL, and Azure ML.
- Mentored junior team members through rigorous code developments and reviews and partnered with cross-functional stakeholders to deliver production-ready analytics solutions using Agile methodologies, CI/CD pipelines, and Azure DevOps.

Data Scientist | Infosys Limited - Client : Energy | India | Aug 2022 – Sep 2023

- Developed an RFM-based customer risk segmentation model for 200K+ customers using SQL, Python for identifying high-, medium-, low- risk debt profiles to enable targeted recovery strategies, increasing recovery rates by 14% and reducing average recovery time by 12 days and implemented model using A/B testing.
- Developed executive-level Power BI dashboards visualising customer segments, debt ageing, recovery performance, and portfolio exposure, enabling data-driven decision-making for senior leadership.

Associate Data Scientist | Infosys Limited - Client : Health | India | Feb 2022 – Aug 2022

- Processed over 1M+ insurance claims and policy records using SQL to identify operational bottlenecks and key drivers of SLA breaches, supporting process improvement initiatives.
- Developed Tableau dashboards and analytics solutions for business reporting by enabling efficient data transformation and feature engineering, extraction.

Academic Projects

AI Job-Search & Application Copilot (Multi-LLM GenAI system): Built an end-to-end LLM platform (Python, TypeScript, Supabase) that ingests 2,000+ live roles daily via GitHub Actions from 9 ATS APIs (Eg: Workday, Greenhouse, Lever, Ashby) and job boards, and ranks them with a cost-aware three-model cascade (DeepSeek, Gemini, OpenAI) using rubric based LLM as a judge scoring, two model consensus and ghost job detection. Shipped a Next.js dashboard with a Kanban application tracker and a Chrome (MV3) copilot that autofills applications from CV and drafts context grounded answers that learn from user edits, 140+ automated tests.

RAG Doc Analytics Assistant for Business Reports: Built a RAG document-intelligence assistant using LangChain, FAISS, Streamlit and Gemini API to ingest multi-format reports, generate embeddings, retrieve evidence and produce context-aware answers with page-level citations. Implemented Vector embeddings, semantic retrieval and vector search, automated evaluation monitoring and unit testing.

Health Insurance ML Pipeline: Built an enterprise-style end-to-end data science and analytics platform using PySpark and Delta Lake, implementing medallion layers, fact and dim tables, star-schema marts, churn, fraud and high-cost-claim classification models, MLflow experiment tracking, batch scoring, data-quality monitoring, scheduling, Docker, pytest, Github Actions CI/CD, Azure Container Apps, Power Apps dashboards.

Consumer Review Intelligence: Developed an NLP sentiment-classification pipeline using 10K+ Amazon Books reviews from Hugging Face. Implemented Bag-of-Words and TF-IDF feature extraction and trained SVM, Naive Bayes and Random Forest models with hyperparameter tuning to identify sentiment inconsistencies and support marketing optimisation.

Education

- MSc Data Science | Coventry University, Coventry | May 2025 – May 2026 | Grade: 76%
- Bachelor of Technology Mechanical Engineering | JNTU University, India | Aug 2017 – Jun 2021 | Grade: 79%

Certifications

- Databricks Certified Generative AI Engineer Associate
- IBM Data Science Professional Certificate
- Google Data Analytics Professional Certificate
- Microsoft Certified: Power BI Data Analyst Associate PL-300
- Microsoft Certified: Azure Fundamentals